Day 56 – Full System Simulation

End-to-end pipeline exercise, AI versus human review, findings, and improvement actions

Objective and scope

Validate the Zecpath AI hiring journey from resume intake through ATS scoring, screening, HR and technical interviews, unified scoring, and final recommendation. This run used synthetic candidate DAY56-SYNTH-001 for a Python Developer role.

Simulated pipeline results

Stage	AI result	What was exercised
Resume intake	Synthetic JSON fixture	Consent checks and fixture loading; no live upload or parser integration.
ATS	95.05 / 100	Skills, experience, education, and semantic match.
Screening	95.00 / 100	Three answers accepted; conversation completed.
HR interview	87.97 / 100	18-question flow completed; scripted Q&A had mismatches.
Technical interview	81.50 / 100	Eight questions; score came from illustrative rubric inputs.
Unified / final	89.39 / 100; Hold / Review	Explicit simulation rule required human review.

Human ratings and comparison

Stage	Human	AI	Human – AI	Assessment
ATS	95	95.05	-0.05	Close agreement.
Screening	95	95.00	+0.00	Exact agreement.
HR interview	90	87.97	+2.03	Provisional: scripted answers were paired with mismatched adaptive questions.
Technical interview	82	81.50	+0.50	Illustrative: human reviewed synthetic sample answers; engine score used rubric inputs.
Final decision	Advance / Recommend	Hold / Review	—	AI hold was triggered by a simulation review gate, not candidate quality alone.

Confirmed inconsistencies

• HR: fixed synthetic answers were fed into adaptive questions in order. For example, the notice-period question received an answer about improving backend skills. The 90 vs 87.97 comparison is provisional.

• Technical: the engine accepted placeholder answer-quality values and did not score free-text. The AI 81.5 used illustrative rubric inputs; the human 82 reviewed synthetic sample answers.

• Integration: screening outputs 0–100 while unified scoring expects 0–10, so the simulation divides by 10. Resume intake used a fixture rather than a real upload flow.

Performance, recommendations, and limits

Performance from one local run

Measure	Observed	Context
Total time	1.8341 seconds	One candidate; cached model; first-time model download excluded.
Slowest stage	ATS scoring: 1.1669 seconds	Includes local model initialization.
Screening	0.3143 seconds	Three-answer synthetic flow.
HR / technical	0.0111 / 0.0002 seconds	Deterministic / placeholder-answer paths.
Legacy screening benchmark	10 / 10 matched (100%)	Intent/action labels only; small existing dataset.

Recommended improvements

• Pair each HR response with the adaptive question it answers, persist the transcript, and reject or clarify off-topic responses before scoring.

• Score actual technical free-text answers against a transparent rubric; retain the questions, answers, and rubric evidence for human review.

• Standardize stage score ranges and validation at module boundaries instead of relying on an implicit divide-by-10 conversion.

• Connect real PDF/DOCX upload and parser output to the ATS schema, then run tests using representative resumes and jobs.

• Expand human evaluation beyond one synthetic candidate and 10 screening intent cases; assess agreement, fairness, and performance across a representative dataset.

Limitations

• This is one local synthetic candidate run. Timings do not establish production throughput, reliability, or fairness.

• Resume intake used an embedded JSON fixture; there is no production file-upload path or parser-to-ATS adapter in this run.

• The HR transcript used synthetic answers assigned sequentially to adaptive questions; several prompt-answer pairs were mismatched, making the AI HR score provisional.

• Technical interview flow was exercised, but its AI score came from illustrative rubric values rather than evaluation of free-text answers. The human rating used a synthetic answer sheet.

• Human review reflects one reviewer who had access to simulation results; the comparison was not blinded or a multi-reviewer validation.

• The separate screening intent/action benchmark covers 10 existing labeled cases only; its 100% agreement does not demonstrate broad accuracy.

Overall status

The synthetic pipeline run and human ratings are documented. ATS and screening scores closely align; HR and technical comparisons are illustrative because their simulated evidence was not evaluated as a coherent, actual candidate interview. The user's final human decision was recorded as Advance / Recommend based on the statement that all human reviews were okay. Production readiness, broad fairness, and real upload integration remain unvalidated.