Day 64 - Internal Review & System Walkthrough

Internal assessment, improvement priorities, and action plan

Review basis

This assessment is based on the synthetic Day 63 pipeline run and a source review of the runner and project landing page. No mentor or external reviewer feedback was received; this is a simulated internal review and should not be represented as feedback from another person.

End-to-end walkthrough

The four synthetic profiles completed the demonstration flow: resume intake and consent, ATS scoring, screening, HR interview, technical interview, machine-test evaluation, unified scoring, eligibility, and final recommendation. The run completed in about 1.06 seconds locally.

Recommendations: Alex Morgan - Selected (91.83); Riley Chen - Selected (82.82); Taylor Brooks - Selected (79.10); Jamie Rivera - Hold / Review (63.09).

Accuracy findings

The run confirms stage connectivity, but it does not measure decision accuracy: the four profiles are synthetic and have no human-verified reference outcomes. HR scores are tightly grouped from 87.93 to 90.12. The limited-evidence profile received an 82.04 screening score, showing that clarity and consistency can outweigh missing job-related evidence.

Technical interview and machine-test scores are illustrative fixtures. Interview answers are not automatically graded, and candidate code is not executed. Resume upload and parser integration are also outside this runner. These recommendations are demonstration outputs and require human review.

Usability findings

The repository scan found a static landing page, not an interactive candidate or recruiter workflow. The command-line walkthrough does not exercise resume upload, score review, interview completion, or error recovery. The landing page hides its navigation on narrow screens without a replacement menu. Usability should be evaluated with actual candidate and recruiter tasks once the interactive workflow is available.

Performance findings

The measured run took about 1.06 seconds for four synthetic profiles. ATS scoring was the largest measured stage, ranging from about 0.04 to 0.56 seconds per profile. These timings come from one small local run and do not establish performance under larger workloads or concurrent use.

Prioritized action plan

1. **Keep human review mandatory.** Treat automated recommendations as decision support; do not use them alone for hiring decisions.
2. **Validate scoring against human-reviewed outcomes.** Build a consented reference set for each role; measure false positives, false negatives, score agreement, and group-level differences before calibrating thresholds.
3. **Replace fixture evaluations.** Grade technical answers against a documented rubric and run coding tasks in an isolated environment using real test cases.
4. **Exercise the complete input path.** Walk through resume upload and parsing, semantic screening, consent handling, and common failure cases.
5. **Review usability.** Test candidate and recruiter tasks with users; provide mobile navigation on the landing page.
6. **Benchmark performance consistently.** Record repeated runs, cold and warm timings, larger candidate sets, and concurrent-use behavior before optimizing.

Conclusion

The project is suitable for a controlled demonstration of pipeline orchestration. Accuracy, usability, and production performance need further validation before the system is relied on for real hiring decisions.

Evidence reviewed: reports/day63_end_to_end_simulation.json; main_day63_end_to_end_simulation.py; website/index.html.