Day 67 - Mock Demo Feedback Report

Review date: 03 November 2026 | Rehearsal format: simulated stakeholder session

This review combines a local run of the synthetic Day 63 hiring-pipeline harness with a role-play Q&A. No external stakeholder or mentor participated, so the feedback below is simulated rather than a record of real reviewer comments.

Rehearsal result

- The current harness processed four fictional candidate profiles across two Python roles and eight configured pipeline stages.
- Two fresh runs produced matching candidate scores. The runner reported 0.60 seconds and 0.35 seconds for the measured pipeline; its timer starts after CandidateMatchingService construction, so model initialization is excluded.
- The first cold start needs a Sentence Transformer preflight. The model metadata request required Hugging Face access in the rehearsal; once cached, the model loaded offline. Warm the model before the audience arrives.
- HR interview scores span only 87.93 to 90.12 across the four profiles. The demo heuristic does not separate profile strength reliably and should not be presented as a validated ranking signal.

Current results to present

Strong synthetic profile 001: ATS 95.59, screening 85.83, technical 87.20, machine-task fixture 92.40, unified score 90.45; simulated disposition: Selected. Limited-evidence profile 004: ATS 70.60, screening 72.87, technical 42.40, machine-task fixture 31.00, unified score 61.71; simulated disposition: Hold / Review.

The saved Day 63 JSON report contains older screening and unified scores (for example, profile 001 screening 95.00 and unified 91.83). Run the demo before presenting so the report is refreshed from the current code; the improved deck uses the fresh rehearsal values.

Mock stakeholder questions

- Does it process uploaded resumes? No; this harness reads structured synthetic JSON.
- Does it execute candidate code? No; machine-test scores are supplied fixtures.
- Can we act on a Selected result? No; it is a simulated disposition, and a human must make real decisions.
- What impact is proven? None yet; a pilot needs a measured baseline and governed outcome review.

Feedback, changes and readiness

The improved presentation keeps the original visual style and adds an explicit run-of-show and stakeholder Q&A.

Weak areas found and fixes made

- Timing was not visible. The deck now sets a six-minute demo and four-minute Q&A, with a timebox and segment notes; it also calls for a model preflight.
- The AI scope could be overstated. The Q&A distinguishes ATS semantic matching from rubric heuristics, and the slides state that uploads, parser integration and a candidate-facing app are not exercised.
- The code task could sound automated. The deck says task scores are fixtures and candidate code is not run; it also labels recruiter review as outside the harness.
- The saved result JSON was stale. The score slide now uses two matching fresh runs and instructs the presenter to refresh the report before the demo.
- No business result is measured. The deck presents impact as a hypothesis and names pilot measures instead of claiming gains.

Six-minute run of show

0:00-0:15 opening; 0:15-0:45 problem; 0:45-1:25 architecture; 1:25-2:10 flow; 2:10-3:10 live run; 3:10-4:00 results; 4:00-4:35 business impact; 4:35-5:05 human review; 5:05-5:50 Q&A preview and close. Keep ten seconds as a buffer; reserve four minutes for questions.

Pre-demo checklist

- Before the session, activate the virtual environment at the repository root, run python main_day63_end_to_end_simulation.py, and confirm the JSON report refreshed.
- Check model availability. If the model is cached and the machine is offline, set $env:HF_HUB_OFFLINE = '1' in PowerShell.
- Use only synthetic profiles; open the refreshed report and keep the four-minute Q&A timebox.

Final demo readiness

READY for a synthetic mock demo after model preflight and report refresh. NOT READY for production hiring decisions: this rehearsal does not validate accuracy, fairness, time savings or real candidate outcomes, and the harness is not an upload-to-decision product.

The Q&A, reviewer perspectives and feedback in this report are simulated. They should be replaced with actual stakeholder feedback when available.