ZECPATH AI - AI Roadmap

Day 69 submission | Forward-looking plan, not a delivery commitment

Vision

Develop ZECPATH AI into a governed decision-support workflow that helps recruiters organize job-related evidence while preserving human control over hiring decisions. Keep this prototype separate from the company official website unless a future integration is explicitly approved.

Starting point

The repository has callable parsers, semantic ATS matching, screening and interview components, unified scoring, recommendations, and compliance-support utilities. Day 56 and Day 63 scripts demonstrate selected synthetic flows. Day 68 passed 679 tests and 8 subtests. The codebase is not a production service: parser output is not automatically wired into ATS matching, there is no implemented HTTP API, and demo technical/machine-test results include supplied rubric values.

Phase 1 - Connect a controlled end-to-end service

Create an orchestration layer that validates resume and job schemas, invokes parsing and matching, checks consent before processing, passes typed results between stages, and returns an explainable recommendation with source evidence. Add clear errors and trace IDs without exposing unnecessary candidate information in logs.

Exit gate: integration tests cover PDF/DOCX input, malformed files, schema edge cases, consent denial/withdrawal, and repeated requests. Reports remain traceable and internally consistent.

Phase 2 - Enforce privacy, access, and retention

Connect consent and access checks to every service entry point. Add authentication and authorization, protected storage, and a retention worker that performs approved deletion rather than only calculating an action. Define data minimization, backups, incident response, and access review.

Exit gate: security review confirms access restrictions; consent and withdrawal are enforced; retention affects storage as specified; logs and exports minimize personal data.

Phase 3 - Validate scoring quality and fairness

Build a governed, consented evaluation set across job families. Obtain independent human ratings, measure reviewer agreement, compare model outputs with the reference, and analyze errors and group-level outcomes before calibrating thresholds. Keep protected attributes out of candidate scoring.

Exit gate: acceptance criteria are defined before testing; results can be reproduced; subgroup limitations are documented; reviewers can inspect evidence and correct or override outputs.

Roadmap - Safe Pilot and Operations

Advance only when the evidence and review gate for each phase is met

Phase 4 - Improve interview and technical evaluation safely

Prioritize evidence-grounded summaries and structured rubric support. Measure transcription quality and candidate experience separately from interview scores. If coding tasks are run, execute untrusted submissions only in an isolated sandbox with resource limits and no access to production data or secrets. Do not present fixture-scored tests as verified code execution. Keep emotion and sensitive-trait inference out of ranking.

Exit gate: scores cite reviewable evidence; outputs are compared with independent human ratings; code cannot access the host; consent and accommodations are handled before a pilot.

Phase 5 - Run a limited internal pilot

Start with a small, approved set of roles and users in a separate internal application or service. Define support ownership, monitoring, reliability and latency goals, reviewer correction tracking, candidate notice, escalation, rollback, and a feature-disable path. Do not connect the official company website until its owners approve the integration and security review.

Exit gate: the pilot meets pre-agreed quality, safety, privacy, fairness, and operational criteria. Expand only the features that pass; revise or pause the rest.

Phase 6 - Scale selectively

Expand to additional roles or stages only after pilot evidence remains strong as data, models, and workloads change. Version models and rubrics, monitor drift and failures, preserve an audit trail, and maintain a practical rollback or disable mechanism.

Roadmap principles

- Keep humans accountable for real employment decisions.
- Explain outputs with job-related evidence, uncertainty, and missing information.
- Keep synthetic demos separate from live candidate records; minimize and protect personal data.
- Re-evaluate after a model, rubric, job family, or input source changes.
- Set measurable targets with stakeholders before pilots; this document invents no dates or performance guarantees.

Relationship to Day 58

This roadmap refines the Day 58 AI feature proposal using the Day 62 module reference and Day 68 validation. Interview coaching and transcript summaries remain proposals until implemented and evaluated against the gates above.